



re-forge

re-forge

The *compounding layer*
for your AI toolchain.

Akasxh & Shine · UCWS Singapore 2026

● Top 5 · 130+ votes

● 2× ICML 2026

● open-source · model-agnostic


✦ re-forge.vercel.app



re-forge.vercel.app



github.com/Akasxh/re-forge



join the waitlist

■ THE HOOK

Every session starts from *zero*.

You run dozens of AI coding sessions a day across Claude Code, Cursor, Codex, ChatGPT. They forget what worked, repeat your mistakes, and leave no trace.

You are the only memory in the loop. That doesn't scale, and it doesn't compound.

0

memory carried between sessions, every tool, every day

1

thing that remembers what worked yesterday: **you**

THE PROBLEM

A single agent is confident,
and *unaccountable*.

01 It hallucinates, anchors on its first idea, and forgets yesterday.

02 On long-horizon work that is wasted tokens, repeated bugs, and **no audit trail**.

You can't prove what it checked.

03 The most valuable thing, **how your team actually works**, evaporates.

Gone at the end of every session.

■ THE INSIGHT

The edge isn't a smarter model.

Everyone rents the same models. Renting a bigger one buys nobody a moat.

The edge is the *operating layer* on top of them.

- coordination
- adversarial verification
- evidence on disk
- memory that compounds

■ WHAT RE-FORGE IS

One prompt → a whole *org.*

A coordinated org of specialist agents (research, engineering, security, testing, docs, evolution) plus a Forge meta-agent. Model-agnostic: it's a procedure, not a model.

- **Specialists fan out** from one lead; each lane works in parallel.
- **Skeptic + adversary attack** the work, so weak claims die before you see them.
- **Survivors reach SYNTHESIS:** evidence lands on disk, memory gains a lesson.



PROOF · RECEIPTS, ALL REAL AND CLICKABLE

Proven on the hardest user there is: *itself*.

01

Compute-Aware Mixture-of-Agents: Verifier-Gated Adaptive Aggregation under a Fixed Token Budget

Shine Gupta · S Akash · SCALE Workshop, ICML 2026

public at camera-ready

02

Who Hallucinates Tools, How Often, and What Fixes It?

S Akash · Shine Gupta · SCALE Workshop, ICML 2026

public at camera-ready

Top 5

· 130+ votes

a top-5 project on the hackathon leaderboard

685

checks

machine checks: green CI on every push

128

skills · 6

teams

it tests itself the way it tests your code

THE DEMO · SECURE JWT AUTH

● PLAIN AGENT · GENUINELY GOOD

A plain agent built JWT auth, and did it well.

19 / 19 its own security tests passed

The token core is sound: Argon2id, refresh rotation, reuse-detection, revocation. re-forge's report says so. We give credit.

By every signal the agent gives you, this looks shippable.

● RE-FORGE SECURITY TEAM · VERIFIED

One **HIGH** hole the 19 tests missed:

- ✗ **Broken object-level authorization (IDOR):** any logged-in user can **read and DELETE** every other user's data.

reproduced cross-user · against the running app

+ MEDIUM GAPS


- ✗ **No rate limiting:** brute-force wide open
- ✗ **Username enumeration:** login reveals who exists

+ 3 more medium

adversary refused to inflate; it doesn't cry wolf

verdict: ADVISORY · 1 HIGH + 5 MEDIUM

Green tests are not secure. *That gap ("looks done" vs "verified") is the product.*



THE PRODUCT

✓ ships today for Claude Code

A memory you *own*.

Today, re-forge already watches your Claude Code sessions and forges skills from them. **The company makes that always-on and cross-tool.**

It sits quietly across your tools, observes how your team actually works, and forges personalized, self-evolving skills & plugins, automatically.

roadmap: always-on cross-tool monitor

THE PRODUCT • POSITIONING


The *discovery layer* for your AI workforce.

The bottleneck isn't the model. It's *discovery*: no one can see how their agents actually work, so improving them is guesswork.

re-forge watches the agents (**not your people**) across every tool, finds what they repeat and what they get wrong, and **forges the skills that fix it**.

Self-evolving. The shared brain your whole team's agents draw from.

■ Product preview • mockup, not live data



THE PRODUCT · DISCOVERY

WORKSPACE

Overview

Discovery12

Skills9

Team Memory

Audit

AE

acme-eng8 seats

WATCHING

Claude Code

Cursor

Codex

ChatGPT

Product preview · mockup · not live data

YV

1,284sessions observed

12patterns surfaced

9skills forged

~18h/wkrepeat-work avoided

12 patterns your agents keep repeating

ranked by repeat-cost · last 30 days

Discovery

Audit

1

Leakage check re-derived by hand

Claude Code ×5

Cursor ×2

seen in

✕ Forge skill → quant-leakage-guard

2

JWT refresh-rotation hand-rolled

Cursor ×4

Claude Code ×3

one shared skill covers both

✕ Forge skill

3

Same OWASP checklist pasted into prompts

ChatGPT ×6

Codex ×3

9 sessions across 4 repo

✕ Forge skill

Dismiss

Forged skills

auto-authored from patterns

quant-leakage-guard

forged 3d ago · used 14× · ▲ adoption3 teammates

jwt-rotation

used 8× · 2 teammates ▲ adoption

owasp-preflight

used 5× · 1 teammate steady

It finds the work your agents repeat, and forges the skill that ends it.

re-forge

🌟 reforge.vercel.app

THE PRODUCT · COMPOUNDING + AUDIT

Forged skills

adoption climbs with use

◆ quant-leakage-guard

forged 3d ago · used 14× · 3 teammates

▲ adoption

◆ jwt-rotation

used 8× · 2 teammates

▲ adoption

◆ owasp-preflight

used 5× · 1 teammate

steady

Team memory

the shared brain

47 lessons · shared across 8 seats

● purge + embargo before any time-series CV

used 31× · top lesson

● evidence-first: every claim links its source

used 26×

AUDIT LOG

project acme/api · 8 seats

Product preview · mockup · not live data

tail · live

#4127

Claude Code 14:02:11

● SECURITY GATE · KILLED

evidence ↗

killed claim *"all inputs are validated"*:
unsanitized field reached the query.

#4119

Cursor 13:48:39

● ADVERSARY · REFUTED

evidence ↗

refuted *"no known CVEs"*: sole source was a
vendor blog post.

#4108

Claude Code 13:31:02

● VERIFIER · PASSED

evidence ↗

confirmed *"Argon2id params meet policy"*:
reproduced & linked.

Open any decision: see who killed it, and why. Nothing is trusted until something tried to kill it.

THE PRODUCT · TEAM MEMORY

One project. One *shared brain*.

Your coworker's agent just made a decision; yours already knows. No duplicated work. No contradictions. The team decides *once*, and every agent respects it.

■ **Product preview** · mockup, not live data

✓ **ships today**: the concurrent file-locked memory merge (flock + atomic-rename, zero lost writes). **roadmap**: the cross-user LIVE sync.

Measured · a real 2-dev parallel merge: the shared spec reconciled both contracts before code, turning **13 conflict hunks** across 4 files into **0. ~25–45 min** of rebase-and-debug saved per overlapping feature pair · modeled. [watch the demo ↗](#)

Siloed today across every tool and teammate. re-forge makes it one synced memory you own.

Live team memory

shared context across teammates

Product preview · mockup

project acme/payments

2 active now A S

SHARED CONTEXT · LIVE FEED

S

Shine's agent · 12m ago · **decided**: Argon2id + refresh rotation

evidence ↗

Akash's agent builds on this, won't re-decide it

A

Akash's agent · 4m ago · **found**: Stripe webhook needs an idempotency key

saved to shared memory

⚠

duplicate work prevented · Shine and Akash both queued "add rate limiter" → merged

merged into one task

SHARED DECISION LOG · every agent on the project respects these

📁 auth = Argon2id + rotation

📁 DB = Postgres, no ORM

📁 no same-day signals in any backtest

■ WHY IT COMPOUNDS


Every session makes it *stickier*.

The longer a team uses it, the more it knows how they work. Your layer gets better with use, the opposite of a tool you outgrow.

- 01 That personal / team memory is **your data, your patterns**.
- 02 No base model and no competitor can replicate it.
- 03 **A real network effect, per team.**

■ MARKET · TAM / SAM / SOM

A large market, entered through a *small, defensible wedge*.



- **TAM · \$52.6B by 2030**
global AI-agents market, up from \$7.8B in 2025 (46.3% CAGR).
- **SAM · ~\$8.1B**
~22.6M AI-agent developers reachable at a typical AI-tool seat price.
- **SOM · ~\$40M ARR**
~0.5% of that base (~110K paid seats). Entry wedge is far smaller: the first 1K to 3K paid seats.

Bottom-up SAM: 36.5M pro developers × 62% on an AI coding agent = ~22.6M, × ~\$360 / seat / yr (assumed industry benchmark) = ~\$8.1B.

SOURCES

MarketsandMarkets

Grand View Research

SlashData 2025

JetBrains 2025

MONETIZATION • OPEN-CORE + ENTERPRISE

PLANNED • PRE-REVENUE

Open-core. Enterprise, *configured per team.*

• FREE • OSS

Free

Individual devs, self-hosters, bottom-up inside teams.
The adoption wedge.

- Open-source operating + **memory** + **skill** layer.
- Local skill-forging from your own Claude Code / Cursor / Codex sessions.
- Personal memory, single user.

where we monetize

• TEAM

Configured per

team

Eng teams standardizing AI agents across tools.
Priced to the value we meter, not a public seat tag.

- **Shared team memory** + cross-session skill forging across the team.
- **Spec-sync**: reconciles parallel work before it becomes a merge conflict.
- Sync across tools, role-based access, usage analytics.

value-based • we configure it with you

• ENTERPRISE

Configured

Large orgs with security, compliance, and on-prem needs.

- **SSO / SAML**, full audit trail of agent activity.
- **On-prem / VPC memory**, data-residency & retention controls.
- SOC 2, admin governance over forged skills.

Land with **one developer's free layer**, expand to the **team's shared brain**, then the **enterprise audit trail** – pricing **configured to each team**, never a public tag.

■ COMPETITION • WHO'S IN THE GAME

Everyone builds the agents or watches them. *We're the layer.*

Single-agent assistants & editors

GitHub Copilot • Cursor • OpenAI Codex • Amazon Q

They generate the agent sessions. re-forge watches across all of them and forges skills + shared memory.

THE SUBSTRATE WE SIT ON

Agent frameworks & orchestration

LangGraph • AutoGen • CrewAI

They build agents. re-forge persists what agents learn: a layer above any framework, not another SDK.

A LAYER ABOVE

Observability & evals

LangSmith • Phoenix • Langfuse • Helicone

They show what happened. re-forge acts on it, turning sessions into reusable skills, memory, and an audit trail.

THE WRITE-BACK LAYER

Agent memory & skill infra

Mem0 • Letta • Zep

Closest adjacency: developer SDKs for app builders. re-forge is a team product that auto-forges skills + shared memory from real coding-agent sessions, with SSO / audit / on-prem.

CLOSEST ADJACENCY

✦ RE-FORGE

The model- and tool-agnostic **memory + skill + audit layer on top**. Everyone else builds the agents or watches them; re-forge is the layer that makes them **learn, remember, and stay auditable**, across every tool.

■ THE VISION

The *compounding layer* for your AI toolchain.

Model-agnostic. Tool-agnostic. The longer your team runs it,
the more it knows how you work, and the harder it is to leave.

● 2× ICML 2026

● Top 5 · 130+ votes

● 685 checks green

Try it now · join the waitlist →



join the
waitlist



re-
forge.vercel.app



github.com/Akash/re-
forge